

TEAM BROTHERHOOD

AgentOps AI

The Enterprise Control Plane for AI Agents

Live on Vercel · Railway · Base Sepolia — a production-grade trust layer for AI agent governance.

THE PROBLEM

Every company has an AI agent sprawl problem — and no one can see it.

Agents pile up across LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, n8n, MCP servers, and custom code — with no shared system of record.

What agents do we actually have?

No unified inventory across frameworks and teams.

What are they doing — and what do they cost?

Token spend, API calls, and compute are invisible.

Who owns each one, and what can it access?

Permissions and ownership are undocumented and unenforced.

What's the highest-leverage thing to fix first?

Without a ranked plan, teams optimize the wrong things.

One control plane, above every framework — not another one to replace them.



Discover

Inventory every agent across frameworks in one place — no manual spreadsheets, no blind spots.



Diagnose

AI-generated Executive Report scoring cost, security, and operational risk with explainable findings.




Act

A ranked Optimization Plan — Immediate Wins, 30-Day, 90-Day, Long-Term Architecture — every recommendation tied to an inspectable rule.

📄 Every plan rating — priority, effort, risk, confidence — comes from an explicit, inspectable rule. Never a black box.

From zero to a full fleet report in minutes.



Connect

Scan

Report

Verify

The nine-section Executive Report includes: executive summary, org overview, cost analysis, security risks, operational risks, optimization opportunities, business impact, priority actions, and an overall health score.


Wallet-Native Auth

OKX Wallet challenge-response sign-in. No passwords. No third-party identity provider. Google, Microsoft, and Okta/SAML are reserved as future sibling providers.

Live Today

A searchable activity timeline and at-a-glance overview — agents found, monthly cost, open risks, and optimization opportunities — available immediately after any scan.

Built as real infrastructure, not a demo.



Trust by Design

Only a hash and minimal metadata ever reach the chain. Report content and org data never leave the database. A chain failure never blocks a scan — the report view simply omits the integrity card.

Pluggable Connectors

GitHub is live today, scanning `requirements.txt`, `pyproject.toml`, and `package.json` for framework markers. LangGraph, CrewAI, MCP, Kubernetes, and cloud-provider connectors are architected in for later phases.

Land with platform teams, expand into the enterprise control plane.

1

Phase 1 — Now

Free, self-serve Health Scan targeting AI-forward engineering teams and Web3-native orgs. GitHub connector: no token, no setup, results in one scan.

2

Phase 2 — Pilots

Enterprise pilots with platform and security teams. Land via a single Health Scan, expand into continuous monitoring for compliance and cost control.


3

Monetization

Tiered plans: Free / Team / Enterprise. Pricing served live from an on-chain contract. PaymentProvider abstraction scaffolded for future x402 metering.

📄 **Defensible moat:** Framework-agnostic by design + a cryptographically verifiable audit trail no competitor currently offers.

What's next for AgentOps Cloud.



Team Brotherhood

A focused team building the infrastructure layer for AI agent governance — six pages, one job: give every enterprise a clear picture of its AI agent fleet and a ranked list of what to do about it.

What we're looking for

- Pilot design partners in platform and security teams
- Seed funding to accelerate connector development
- Hackathon judging criteria and ecosystem partnerships